

April 23, 2020

Final Report

1 Setting up

1.1 Change work directory

```
[1]: import os
os.chdir("/Users/WaferHsu/Desktop/AU/2020Spring/STAT-696-NLP/
↪FinalProj_Alex_Wafer")
```

1.2 Import packages

```
[2]: # nltk
import nltk
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
stopWords = set(stopwords.words('english'))

# numpy
import numpy as np

# regex
import re

# pandas df
import pandas as pd

# plot
import matplotlib.pyplot as plt
import seaborn as sns

# used to remove Pandas warnings
```

```

import warnings
warnings.filterwarnings('ignore')

# sentiment analysis (vader)
from nltk.sentiment.vader import SentimentIntensityAnalyzer

# normalize data
from sklearn.preprocessing import MinMaxScaler

# train/test split
from sklearn.model_selection import train_test_split

# Regression
import statsmodels.formula.api as sm
from sklearn import linear_model

# tfidf
from sklearn.feature_extraction.text import TfidfVectorizer

# NMF (topic modeling)
from sklearn.decomposition import NMF

# KMeans Clustering
from sklearn.cluster import KMeans

# classifying modules
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import RidgeClassifier

# sklearn (classification)
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

# Visualizing Clusters
import matplotlib as mpl
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.manifold import MDS

```

1.3 Load original dataset

```
[3]: data = pd.read_csv("../data/kaggle_RC_2019-05.csv", encoding = "utf8")
```

- Take a brief look at the data.

```
[4]: data.head(3)
```

```
[4]:      subreddit      body \
0  gameofthrones  Your submission has been automatically removed...
1      aww  Dont squeeze her with you massive hand, you me...
2      gaming  It's pretty well known and it was a paid produ...

      controversiality  score
0          0          1
1          0         19
2          0          3
```

- Check missing values: Fortunately, we do not need to deal with NA values data.

```
[5]: data.isna().describe()
```

```
[5]:      subreddit      body controversiality      score
count    1000000  1000000          1000000  1000000
unique         1         1              1         1
top         False      False            False      False
freq    1000000  1000000          1000000  1000000
```

- Check the “body” variable is a string variable.

```
[6]: data.dtypes
```

```
[6]: subreddit      object
body      object
controversiality  int64
score      int64
dtype: object
```

```
[7]: type(data["body"][0])
```

```
[7]: str
```

2 Exploratory Data Analysis (EDA)

- Descriptive statistics for continuous variables. The maximum and minimum has a huge difference.

```
[43]: data.describe()
```

```
[43]:      controversiality      score
count    1000000.000000  1000000.000000
mean         0.029583     11.510103
std         0.169434    149.671560
min         0.000000   -889.000000
25%         0.000000     1.000000
50%         0.000000     2.000000
```

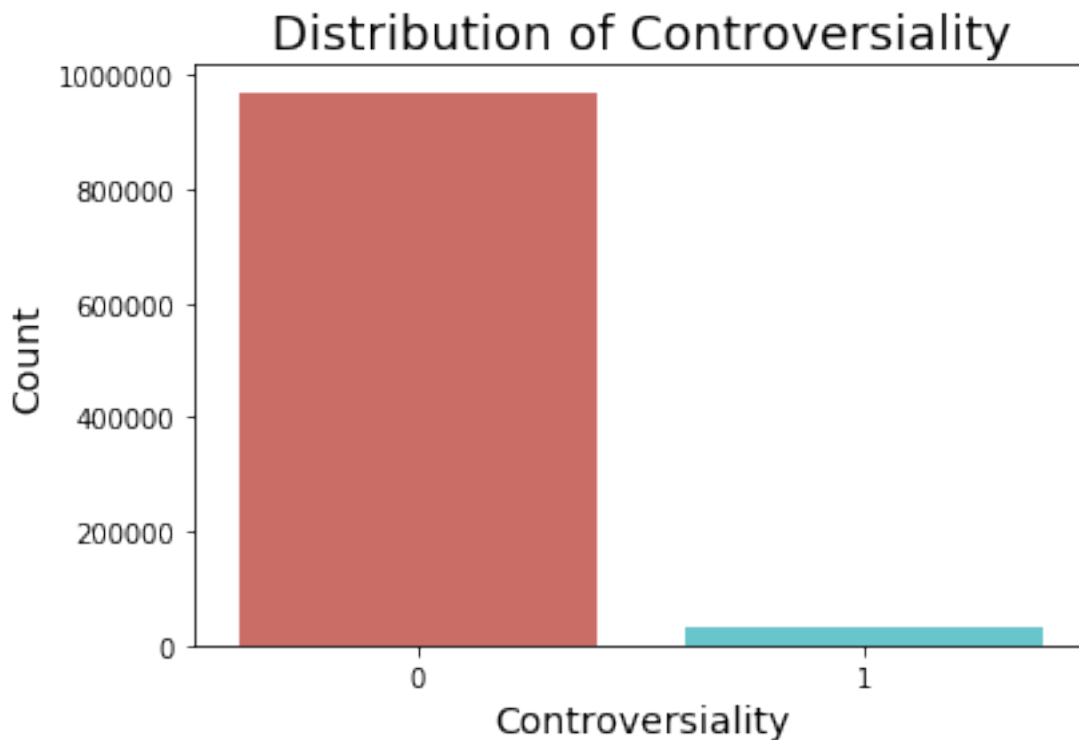
75%	0.000000	4.000000
max	1.000000	35619.000000

2.1 Controversiality

- The distribution of controversiality. It is heavily discrete distributed. There are so many less data in controversiality = 1.

```
[24]: plt.figure(figsize = (6, 4))
sns.countplot(data.controversiality, palette = "hls")
plt.title("Distribution of Controversiality", fontsize = 18)
plt.ylabel("Count", fontsize = 14)
plt.xlabel("Controversiality", fontsize = 14)
```

```
[24]: Text(0.5, 0, 'Controversiality')
```



- The frequency of controversiality.

```
[12]: print(data.controversiality.value_counts())
```


```
0    970417
1     29583
Name: controversiality, dtype: int64
```

2.2 Score

Reddits scores (up reddits minus down reddits) are heavily skewed right. Data mainly distributed lower than 3000, and there are a lot of outliers.


```
[36]: plt.figure(figsize = (18, 3.5))
sns.distplot(data.score, hist = False, color = "red", bins = 10)
plt.title("Distribution of Reddit Scores", fontsize = 18)
plt.ylabel("Frequency", fontsize = 14)
plt.xlabel("Score", fontsize = 14)
```

```
[36]: Text(0.5, 0, 'Score')
```



```
[41]: plt.figure(figsize = (18, 3.5))
data.score.plot(color = 'orange', alpha = 0.8)
plt.title("Reddit Scores for Each Review", fontsize = 18)
plt.ylabel("Score", fontsize = 14)
plt.xlabel("Reddits Reviews (from 0 to 1000000)", fontsize = 14)
```

```
[41]: Text(0.5, 0, 'Reddits Reviews (from 0 to 1000000)')
```



2.3 Subreddit

Frequency Counts of the column “subreddit”. Each category has 25000 reviews. Based on the subreddits, we could sort them out to some broader topics like “sports”, “news & politics”, “movies”, “slangs”, “online games/video games”, and see the classification between two (or more) similar groups. After filtering similar topics out, we are looking at the body text for sentiment analysis,

topic modeling, and clustering to see if there is anything interesting results within the data. This project will focus on the category of politics.

```
[44]: data["subreddit"].value_counts()
```

```
[44]: The_Donald      25000
      gonewild      25000
      Market76      25000
      worldnews     25000
      ChapoTrapHouse 25000
      apexlegends   25000
      marvelstudios 25000
      Animemes      25000
      RoastMe       25000
      soccer        25000
      unpopularopinion 25000
      nfl           25000
      todayilearned 25000
      Pikabu        25000
      gaming        25000
      funny         25000
      videos        25000
      leagueoflegends 25000
      dankmemes     25000
      teenagers     25000
      gameofthrones 25000
      wallstreetbets 25000
      asoiaf        25000
      hockey        25000
      MortalKombat   25000
      memes         25000
      freefolk      25000
      relationship_advice 25000
      trashy        25000
      politics      25000
      movies        25000
      FortNiteBR     25000
      pics          25000
      AskReddit      25000
      aww           25000
      nba           25000
      Showerthoughts 25000
      news          25000
      SquaredCircle  25000
      AmItheAsshole  25000
      Name: subreddit, dtype: int64
```

3 Subset politics data by subreddit and the exploration

```
[46]: politics_data = data[data["subreddit"] == "politics"]
```

```
[47]: politics_data = politics_data.reset_index(drop = True) # reset index
```

```
[48]: politics_data.head()
```

```
[48]:  subreddit                                body \
0  politics  Yes, there is a difference between gentle supp...
1  politics  He also got married, and they filed jointly fo...
2  politics  So you think we can just tell people they no l...
3  politics      ITT, lots of people without jobs complaining.
4  politics                                "You boys wanna shovel some coal?"

   controversiality  score
0                  0      1
1                  0     12
2                  0      1
3                  0     -6
4                  0     17
```


3.1 Distribution of Subreddit

Re-formatted histogram of politics subreddit scores with x-axis limits from (-50, 50) and y-axis limits from (0, 2000).

```
[50]: %%time
sns.distplot(politics_data["score"], kde = False, color = "red", bins = 3000)
plt.title("Distribution of Politics Subreddit Scores", fontsize = 18)
plt.ylabel("Frequency", fontsize = 14)
plt.xlabel("Score", fontsize = 14)
plt.xlim(-50,50)
plt.ylim(0, 2000)
# most scores fall within 0 to 20 and there are a lot more observations between
↪ 0 and 20 than 2000
```

```
CPU times: user 2.67 s, sys: 111 ms, total: 2.78 s
Wall time: 2.79 s
```

```
[50]: (0, 2000)
```



3.2 Creation of New body_length Variable

Counting number of characters in (uncleaned) body variable.

```
[63]: politics_data["body_length"] = [len(el) for el in politics_data["body"]]
```

```
[64]: politics_data.head()
```

```
[64]:  subreddit                                body \
0  politics  Yes, there is a difference between gentle supp...
1  politics  He also got married, and they filed jointly fo...
2  politics  So you think we can just tell people they no l...
3  politics          ITT, lots of people without jobs complaining.
4  politics                                "You boys wanna shovel some coal?"
```


	controversiality	score	body_length
0	0	1	101
1	0	12	145
2	0	1	89
3	0	-6	45
4	0	17	34

3.3 Distribution of body_length (Histogram)

Data visualizing the body_length variable to see the data pattern. The body length is skewed to the right.

```
[62]: sns.distplot(politics_data.body_length, kde = False, color = "blue", bins = 1000)
plt.title("Distribution of Comment Length", fontsize = 18)
plt.ylabel("Frequency", fontsize = 14)
plt.xlabel("Length", fontsize = 14)
plt.xlim(0, 500)
```

[62]: (0, 500)



3.4 Summary statistics on continuous variables

The maximum in score is 9715 but the mean is 13, data is strongly discreated. Same situation in body_length, which maximum is 9156 but its mean is 212.

```
[60]: politics_data.describe()
```

```
[60]:
```

	controversiality	score	body_length
count	25000.00000	25000.000000	25000.000000
mean	0.03224	13.809040	212.050040

std	0.17664	141.260591	298.831188
min	0.00000	-383.000000	10.000000
25%	0.00000	1.000000	59.000000
50%	0.00000	2.000000	113.000000
75%	0.00000	6.000000	242.000000
max	1.00000	9715.000000	9156.000000

4 Pre-processing politics_data

4.1 Spelling error checking

Test abbreviations or spelling errors. To be honest, there are a lot of things to do to fix the spelling error, my code below is incompleted, but this methodology helps me to understand how much more things I need to consider in the future data analysis.

Reference: <https://medium.com/@indreshbhattacharyya/remaking-of-shortened-sms-tweet-post-slangs-and-word-contraction-into-sentences-nlp-7bd1bbc6fcff>

- Test abbreviations: “wd”, “bt”, “der”, “hs”, “dt”, “nt”. According to the output, ‘hs’ = spelling error; ‘dt’ = (probably) doubt; ‘nt’ = not.

```
[65]: abbre_list = ["wd", "bt", "der", "hs", "dt", "nt"]

# make a list by "matchlist", find matches from the "politics_data"
matchlist_test = [" " + abbre_list[i] + " " for i in range(len(abbre_list))]
print("Add space before and after in order to test the word in reddit",
      matchlist_test)
print()

# print all matches by the "politics_data"
print("Matching reddit: ")
print()
for abbr in matchlist_test:
    for rv in politics_data["body"]:
        if abbr in rv:
            print("-----"+abbr+"-----")
            print(rv)
            print()
```

Add space before and after in order to test the word in reddit [' wd ', ' bt ',
' der ', ' hs ', ' dt ', ' nt ']

Matching reddit:

----- hs -----

I think my favourite moment is the one where cartman pretends his hand is a
conman pretending to be Jennifer lopez. The moment Kyle decides to believe him

and he goes Ha ha hs ha ha ha, I got you kinda

----- dt -----

I'm not going to hold my breath. I just listened to Grassley on C-span and they're entrenched into the myth that barr/mueller exonerated dt - and they're continuing to spread that myth/propaganda widely. Today's hearing with barr will be fascinating!

<https://www.youtube.com/watch?v=CK2GNiXdE3c>

----- nt -----

Didnt not nt mean it.

4.2 Text Clean, Tokenize, and Lemmatization

- Write functions to clean and lemmatize “body” variable.

```
[67]: ### text cleaning
def textCleaner(text):
    output = text

    #lowercase
    output = output.lower()

    # fix contractions
    output = output.replace("'", "")
    output = output.replace("can't", "can not")
    output = output.replace("won't", "will not")
    output = output.replace("n't", " not")
    output = output.replace("\'ll", " will")
    output = output.replace("\'ve", " have")
    output = output.replace("\'d", " would")
    output = output.replace("\'s", " is")
    output = output.replace("\'m", " am")
    output = output.replace("\'re", " are")

    # remove stop words
    clean_sentence = []
    for word in output.split():
        if word not in stopWords:
            clean_sentence.append(word)
    output = " ".join(clean_sentence)

    # remove email addresses & Internet domains
    output = re.sub(r'r\/\S+', '', output)
    output = re.sub(r'http\S+', '', output)
    output = re.sub("\S*\@\S*\s?", "", output)
    output = re.sub("\S*\.\edu|\.\com|\.\gov|\.\net\S*\s?", "", output)
```

```

output = re.sub("@\S*", "", output)

# remove punctuations
output = re.sub("[^a-zA-Z\s]", "", output)

return(output)

### tokenize & lemmatization
def textNormalization(corpus):
    output = corpus

    # in the text, lemmatize each word in the tokenized list
    output = [lemmatizer.lemmatize(w) for w in word_tokenize(output)]

    # rejoin all the words into one string by a space
    output = " ".join(output)

    return(output)

```

- Apply functions to the body variable.

```

[82]: politics_data["body"] = politics_data["body"].apply(textCleaner)
      politics_data["body"] = politics_data["body"].apply(textNormalization)

      # briefly print the variable
      politics_data["body"].head()

```

```

[82]: 0    yes difference gentle suppression hard suppres...
      1    also got married filed jointly husband income ...
      2           think tell people longer right express
      3           itt lot people without job complaining
      4                   boy wan na shovel coal
      Name: body, dtype: object

```

4.3 Distribution of *cleaned* body_length (Histogram)

- Change body_length from characters to number of words.

```

[74]: politics_data["body_length"] = [len(el.split()) for el in politics_data["body"]]

```

- The distribution of the body_length variable.


```

[79]: sns.distplot(politics_data["body_length"], kde = False, color = "blue", bins = 1000)
      plt.title("Distribution of Comment Length (Words)", fontsize = 18)
      plt.ylabel("Frequency", fontsize = 14)
      plt.xlabel("Numver of Words", fontsize = 14)

```

```
plt.xlim(0, 80)
```

```
[79]: (0, 80)
```



4.4 Make a Text Corpus

- For body variable, make a corpus for topic modelling.

```
[83]: corpus = list(politics_data["body"])

# briefly print the corpus
corpus[0:7]
```

```
[83]: ['yes difference gentle suppression hard suppression neither good thing',
      'also got married filed jointly husband income teacher number number',
      'think tell people longer right express',
      'itt lot people without job complaining',
      'boy wan na shovel coal',
      'everything power make sure biden get nomination fall line doe primary
      coronation',
      'according mueller bar lied misrepresented pretty much everything democrat need
      impeach ag']
```

5 Sentiment Analysis

5.1 Harvard IV-4 scoring

- Read in dataset, separate to positive and negative values, and drop NA values.

```
[86]: # read Harvard IV-4 sentiment data
harvard = pd.read_excel("./data/dict_Harvard.xls")

# extract positive & negative lists
harvard_pos = harvard[["Entry", "Positiv"]].dropna()
harvard_neg = harvard[["Entry", "Negativ"]].dropna()
```

- Check the lengths of the datasets.

```
[87]: print("Length of positive dictionary: ", len(harvard_pos))    # 1915 rows
      print("Length of negative dictionary: ", len(harvard_neg))    # 2291 rows
```

Length of positive dictionary: 1915

Length of negative dictionary: 2291

- Take a look at the dictionary, all words are capitalized.

```
[92]: print(harvard_pos.head())
      print(harvard_neg.head())
```

```
      Entry  Positiv
8    ABIDE  Positiv
9  ABILITY  Positiv
11   ABLE  Positiv
18  ABOUND  Positiv
41 ABSOLVE  Positiv
      Entry  Negativ
2    ABANDON  Negativ
3 ABANDONMENT  Negativ
4     ABATE  Negativ
6  ABDICATE  Negativ
7    ABHOR  Negativ
```

- Make a function to clean dictionaries (lowercase, remove punctuation).

```
[93]: def harvardCleaner(word):
      output = word

      # lowercase
      output = str(word).lower()
      # remove punctuations
      output = re.sub("[^a-zA-Z]", "", output)

      return(output)
```

- Make a Harvard IV-4 scorer function for applying it to the `politics_data`.

```
[94]: def harvardScorer(text_input, sentiment_list):

    # GET NUMERATOR
    numerator = 0
    for word in str(text_input).split():
        if word in list(sentiment_list):
            numerator += 1
        else:
            numerator += 0
    # GET DENOMINATOR
    denominator = len(str(text_input).split())

    # CALCULATE SENTIMENT SCORE
    if denominator == 0:
        sentiment_score = 0
    else:
        sentiment_score = numerator/denominator

    return(float(sentiment_score))
```

- Apply the function for cleaning dictionaries and drop duplicated words.

```
[95]: # apply function: harvardCleaner
harvard_pos["Entry"] = harvard_pos["Entry"].apply(harvardCleaner)
harvard_neg["Entry"] = harvard_neg["Entry"].apply(harvardCleaner)

# drop duplicated words
harvard_pos.drop_duplicates("Entry", "first", inplace = True)
harvard_neg.drop_duplicates("Entry", "first", inplace = True)
```

- View the length of the dictionaries.

```
[96]: print("Length of positive dictionary: ", len(harvard_pos))    # 1636 rows
      print("Length of negative dictionary: ", len(harvard_neg))    # 2005 rows
```

```
Length of positive dictionary: 1636
Length of negative dictionary: 2005
```

- Create new variables in the data frame to store Positive_Score and Negative_Score.

```
[97]: # Positive Score
politics_data["Positive_Score"] = politics_data.apply(lambda x:
                                                    harvardScorer(x['body'],
↪harvard_pos["Entry"]),
                                                    axis = 1)

# negative score
politics_data["Negative_Score"] = politics_data.apply(lambda x:
```

```

harvardScorer(x['body'],
↳harvard_neg["Entry"]),
axis = 1)

```

- Create new variable “net_score” that is the difference of the positive and negative scores.

```

[98]: politics_data["net_score"] = politics_data["Positive_Score"] -
↳politics_data["Negative_Score"]

```

- Preview the politics_data

```

[99]: politics_data.head(5)

```

```

[99]: subreddit                                body \
0  politics  yes difference gentle suppression hard suppres...
1  politics  also got married filed jointly husband income ...
2  politics                                think tell people longer right express
3  politics                                itt lot people without job complaining
4  politics                                boy wan na shovel coal

   controversiality  score  body_length  Positive_Score  Negative_Score \
0                  0      1           9           0.222222           0.333333
1                  0     12          10           0.100000           0.000000
2                  0      1           7           0.166667           0.000000
3                  0     -6           6           0.000000           0.000000
4                  0     17           6           0.000000           0.000000

   net_score
0  -0.111111
1   0.100000
2   0.166667
3   0.000000
4   0.000000

```

5.2 Slang dictionary scoring

- Import the slang dictionary as a pandas df and change column names.

```

[100]: slangSD = pd.read_csv("./data/dict_slangSD.txt", sep = "\t")
slangSD.columns = ["word", "score"]

```

- There are 96460 rows from slang dictionary, and the word entries have not cleaned yet.

```

[102]: print("Length of dataframe ", len(slangSD))
slangSD.head()

```

```

Length of dataframe  96460

```



```
[102]:      word  score
0      a'f      1
1    a'ight     -1
2  a'nnesia     -1
3    a'pcha      0
4      a's       1
```

- Create a cleaning function for cleaning the slangSD data frame. The dataframe as multi-gram words, punctuations, conjunctions, and different cases that need to be cleaned. This reduces the scoring ability, but is necessary to generate the scores.

```
[103]: def slangCleaner(text):
    output = text

    # fix contractions -1
    contraction_map = {"ain't": "is not", "can't": "cannot",
                       "dont": "do not", "how'd": "how did",
                       "i'd've": "i would have", "let's": "let us",
                       "ma'am": "madam", "o'clock": "of the clock",
                       "shan't": "shall not", "sha'n't": "shall not",
                       "should've": "should have", "so's": "so as",
                       "where'd": "where did", "won't": "will not",
                       "wouldn't": "would not", "y'all": "you all",
                       "you'd": "you would"}
    for i in range(len(contraction_map)):
        output = output.replace(list(contraction_map.items())[i][0],
                                list(contraction_map.items())[i][1])

    # fix contraction - 2
    output = output.replace("n't", " not")
    output = output.replace("'ll", " will")
    output = output.replace("'ve", " have")
    output = output.replace("'d", " would")
    output = output.replace("'s", " is")
    output = output.replace("'m", " am")
    output = output.replace("'re", " are")

    # remove email addresses & Internet domains
    output = re.sub(".*@.*", "", output)
    output = re.sub("\S*\.(com|net)\S*", "", output)

    # remove punctuations
    output = re.sub("[^a-zA-Z\s]", "", output)

    return(output)
```

- Apply clean function, drop duplicates, and preview the slangSD dataset.

```
[106]: # apply function
slangSD["word"] = slangSD["word"].apply(slangCleaner)

# drop duplicates, "first", keeps the first entry
slangSD.drop_duplicates("word", "first", inplace = True)

# reset index
slangSD = slangSD.reset_index(drop = True)

# preview cleaned slangSD
slangSD.head()
```

```
[106]:      word  score
0      af      1
1    aight     -1
2  annesia     -1
3    apcha      0
4     a is      1
```

- Briefly test slangSD to make sure the cleaning function does work. The output looks qualified.

```
[109]: # test slangSD
def text_wd(word):
    for el in slangSD["word"]:
        if word in el:
            print(el)
text_wd("should have")

# 2nd way for testing
slangSD[slangSD["word"] == "have"]
```

```
[109]:      word  score  gram
42646  have      0      1
```

- Test if there is any duplicated word.

```
[110]: slangSD[slangSD["word"] == "a"]
```

```
[110]:      word  score  gram
7      a      0      1
```

- Check the length of the cleaned dataset, there are 94292 rows.

```
[107]: len(slangSD["word"])
```

```
[107]: 94292
```

- Make a new variable that is the count of the words in the slang term -> More than 1-gram words are removed because it will take too much computing power to count sentiments of all

1, 2, 3, etc. gram words and the frequency of multi-gram slang decreases with gram length.

```
[ ]: # keep 1-gram only
slangSD["gram"] = 0

for i in range(len(slangSD)):
    slangSD["gram"][i] = len(slangSD["word"][i].split())

slangSD = slangSD[slangSD["gram"] < 2]
```

- Check the length of the new subsetting data frame – it is reduced by ~ 30,000.

```
[43]: len(slangSD)    # 59,227
```

```
[43]: 59228
```

- Create a slang sentiment score function with nuanced numerator calculation changes.

```
[115]: def slangScorer(text_input, sentiment_list):

    # GET NUMERATOR
    numerator = 0
    for word in str(text_input).split():
        if word in list(sentiment_list):
            numerator += int(slangSD["score"][slangSD["word"] == word])
        else:
            numerator += 0
    # GET DENOMINATOR
    denominator = len(str(text_input).split())

    # CALCULATE SENTIMENT SCORE
    if denominator == 0:
        sentiment_score = 0
    else:
        sentiment_score = numerator/denominator

    return(sentiment_score)
```

- Run slang scorer on data frame and store in new variable — **takes a LONG time!** You may import the .csv for saving time.

```
[46]: %%time
politics_data["slang_score"] = politics_data.apply(lambda x:
    ↪slangScorer(x['body'],
    ↪slangSD["word"]),
    axis = 1)
politics_data
```

CPU times: user 36min 40s, sys: 6.33 s, total: 36min 46s
Wall time: 36min 59s

```
[46]: subreddit                                body \
0      politics yes difference gentle suppression hard suppres...
1      politics also got married filed jointly husband income ...
2      politics think tell people longer right express themselves
3      politics                itt lot people without job complaining
4      politics                you boy wan na shovel coal
...
24995 politics everyone care truth justice incensed no includ...
24996 politics big question senate never openly refused enter...
24997 politics report literally say president cooperative eno...
24998 politics ok let play unfun little game barr gave series...
24999 politics                yay solves problem

      controversiality  score  body_length  Positive_Score  Negative_Score \
0                0      1          9      0.222222      0.333333
1                0     12         10      0.100000      0.000000
2                0      1          7      0.142857      0.000000
3                0     -6          6      0.000000      0.000000
4                0     17          6      0.000000      0.000000
...
24995                0      3          8      0.375000      0.000000
24996                0      2        152      0.105263      0.046053
24997                0      2         14      0.214286      0.071429
24998                0      7         21      0.190476      0.142857
24999                0      0          3      0.000000      0.333333

      net_score  slang_score
0      -0.111111  -0.111111
1       0.100000  -0.200000
2       0.142857  -0.285714
3       0.000000  -0.333333
4       0.000000  -0.500000
...
24995       0.375000  -0.125000
24996       0.059211  -0.125000
24997       0.142857  -0.071429
24998       0.047619  -0.047619
24999      -0.333333   0.333333
```

[25000 rows x 9 columns]

- (optional) Download the “politics_data_slang.csv” for saving the time running the slangSD apply function.

```
[116]: # politics_data = pd.read_csv("../data/politics_data_slang.csv", encoding='utf8')
```

5.3 ANEW dictionary scoring

- Read in the ANEW dataset and change the variable name of the words.

```
[118]: anew = pd.read_csv("../data/dict_ANEW.csv")

anew = anew.rename(columns = {"b'Description": "word"})
anew.head()
```

```
[118]:
```

	word	Word No.	Valence Mean	Valence SD	Arousal Mean	Arousal SD	\
	grin	773.0	7.40	1.87	5.27	2.64	
	honest	210.0	7.70	1.43	5.32	1.92	
	gripe	774.0	3.14	1.56	5.00	2.19	
	honey	792.0	6.73	1.70	4.51	2.25	
	guillotine	196.0	2.48	2.11	6.56	2.54	

Dominance Mean	Dominance SD	Word Frequency
6.00	1.86	13
6.24	2.13	47
4.67	1.79	.
5.44	1.47	25
4.64	2.63	.

- Make a new function to calculate the ANEW scores and append to a new variable.

```
[119]: def anewScorer(text_input, sentiment_list):

    # GET NUMERATOR
    numerator = 0
    for word in str(text_input).split():
        if word in list(sentiment_list):
            numerator += float(anew["Valence Mean"][anew["word"] == word])
        else:
            numerator += 0

    # GET DENOMINATOR
    denominator = len(str(text_input).split())

    # CALCULATE SENTIMENT SCORE
    if denominator == 0:
        sentiment_score = 0
    else:
        sentiment_score = numerator/denominator

    return(sentiment_score)
```

- Run the function through the data set.

```
[120]: politics_data["anew_score"] = politics_data.apply(lambda x:
                                                    anewScorer(x['body'],
                                                            anew["word"]),
                                                    axis = 1)
```

5.4 VADER scoring

- Use the nltk package, VADER is designed with a focus on social media texts. Make a new function to calculate the VADER scores and append to a new variable.

```
[122]: from nltk.sentiment.vader import SentimentIntensityAnalyzer

def vaderScorer(text_input):
    vader = SentimentIntensityAnalyzer()
    output = vader.polarity_scores(str(text_input))['compound']
    return(output)
```

- Apply the vader scorer function to a new variable vader_score.

```
[124]: politics_data["vader_score"] = politics_data["body"].apply(vaderScorer)
```

- Preview data

```
[125]: politics_data.head(3)
```

```
[125]: subreddit                                body \
0  politics  yes difference gentle suppression hard suppres...
1  politics  also got married filed jointly husband income ...
2  politics  think tell people longer right express themselves

   controversiality  score  body_length  Positive_Score  Negative_Score \
0                  0      1           9         0.222222         0.333333
1                  0     12          10         0.100000         0.000000
2                  0      1           7         0.142857         0.000000

   net_score  slang_score  anew_score  vader_score
0 -0.111111  -0.111111    2.222222    0.4203
1  0.100000  -0.200000    0.568000    0.1531
2  0.142857  -0.285714    1.047143    0.0000
```

5.5 Normalize the relevant score variables

- Check the summary statistics of the data – the net_score, slang_score, and anew_score variables need to be normalized so that they can be averaged on the same scale.

```
[126]: politics_data.describe()
```

```
[126]:
```

	controversiality	score	body_length	Positive_Score \
count	25000.00000	25000.000000	25000.000000	25000.000000
mean	0.03224	13.809040	18.679800	0.106885
std	0.17664	141.260591	24.571492	0.116847
min	0.00000	-383.000000	0.000000	0.000000
25%	0.00000	1.000000	6.000000	0.000000
50%	0.00000	2.000000	11.000000	0.090909
75%	0.00000	6.000000	22.000000	0.166667
max	1.00000	9715.000000	660.000000	1.000000

	Negative_Score	net_score	slang_score	anew_score	vader_score
count	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000
mean	0.105858	0.001027	-0.053783	0.568576	0.005949
std	0.116204	0.164613	0.146260	0.690321	0.514391
min	0.000000	-1.000000	-1.000000	0.000000	-0.998900
25%	0.000000	-0.071429	-0.121951	0.000000	-0.401900
50%	0.084826	0.000000	-0.015625	0.402500	0.000000
75%	0.166667	0.076923	0.000000	0.860808	0.421500
max	1.000000	1.000000	1.000000	8.100000	0.994300

- Standardize variables: Import MinMaxScaler which subtracts the minimum value from each score and divides by the range. This rescales each score from a 0-1 scale. They can then be compared with 0 indicating low valence and 1 indicating high valence.

Reference: <https://towardsdatascience.com/scale-standardize-or-normalize-with-scikit-learn-6ccc7d176a02> <https://stackoverflow.com/questions/24645153/pandas-dataframe-columns-scaling-with-sklearn>

```
[127]: from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
politics_data[["net_score",
               "slang_score",
               "anew_score",
               "vader_score"]] = scaler.
↳fit_transform(politics_data[["net_score",
                               ↳
↳"slang_score",
                               ↳
↳"anew_score",
                               ↳
↳"vader_score"]])
```

- Take the average score and store in a new overall score variable which will be used in the regression.

```
[128]: politics_data["overall_sent_score"] = (politics_data["net_score"] +
                                              politics_data["anew_score"] +
```

```
politics_data["slang_score"] +
politics_data["vader_score"]) / 4
```

- Briefly view the dataset.

```
[131]: print("Length of data frame: ", len(politics_data))
politics_data.head(3)
```

Length of data frame: 25000

```
[131]: subreddit                                body \
0  politics  yes difference gentle suppression hard suppres...
1  politics  also got married filed jointly husband income ...
2  politics  think tell people longer right express themselves

   controversiality  score  body_length  Positive_Score  Negative_Score \
0                  0      1           9         0.222222         0.333333
1                  0     12          10         0.100000         0.000000
2                  0      1           7         0.142857         0.000000

   net_score  slang_score  anew_score  vader_score  overall_sent_score
0  0.444444    0.444444    0.274348    0.712021         0.468815
1  0.550000    0.400000    0.070123    0.577965         0.399522
2  0.571429    0.357143    0.129277    0.501154         0.389751
```

- Descriptive statistics for overall_sent_score.

```
[132]: politics_data["overall_sent_score"].describe()
```

```
[132]: count      25000.000000
mean           0.386989
std            0.082609
min            0.120641
25%            0.323939
50%            0.381692
75%            0.448989
max            0.814123
Name: overall_sent_score, dtype: float64
```

6 Regression

6.1 Train Multiple Regression on Pandas DF

- Make controversiality a factor variable

```
[326]: politics_data['controversiality'] = politics_data['controversiality'].
        ↪astype('category')
```


- Find and remove outliers

```
[327]: # Find IQR
Q1 = politics_data["score"].quantile(0.25)
Q3 = politics_data["score"].quantile(0.75)
IQR = Q3 - Q1

# Remove outliers
politics_out = politics_data[politics_data["score"] < (Q3 + 1.5 * IQR)]
politics_out = politics_out[politics_out["score"] > (Q1 - 1.5 * IQR)] #
↳ 21461 rows

# Reset Index
politics_out = politics_out.reset_index()
```

- Train/test split

```
[328]: # Make new df with dependent variables
politics_reg = politics_out[["overall_sent_score", "controversiality",
↳ "body_length"]]
```

```
[329]: from sklearn.model_selection import train_test_split

# subset y variable
y = politics_out.score

# split train and test data
X_train, X_test, Y_train, Y_test = train_test_split(politics_reg, y,
↳ test_size=0.25, random_state = 1)
```

- Fit the model

Reference: <https://towardsdatascience.com/train-test-split-and-cross-validation-in-python-80b61beca4b6>

```
[330]: from sklearn import linear_model
lm = linear_model.LinearRegression()

model = lm.fit(X_train, Y_train)
predictions = lm.predict(X_test)
```

```
[332]: # create a df
df = pd.DataFrame({'Actual': Y_test, 'Predicted': predictions})
df = df.reset_index(drop = True)


# plotting
plt.figure(figsize = (16, 4))
df.plot(kind = 'bar', figsize = (16,5))
```

```
plt.grid(which = 'major', linestyle = '-', linewidth = '0.5', color = 'green')
plt.grid(which = 'minor', linestyle = ':', linewidth = '0.5', color = 'black')
plt.xlim(1400,1430)
plt.title("Comparison between y_true & y_pred", fontsize = 18)
plt.ylabel("Value", fontsize = 14)
plt.xlabel("Order of Entries", fontsize = 14)

# take a look at part of the data
plt.xlim(1920,1950)
```


[332]: (1920, 1950)

<Figure size 1152x288 with 0 Axes>



```
[139]: plt.scatter(Y_test, predictions)
plt.xlabel("True Values")
plt.ylabel("Predictions")
```

[139]: Text(0, 0.5, 'Predictions')



```
[140]: print("R-squared:", model.score(X_test, Y_test)) # low value
```

R-squared: 0.024614384673189327

```
[141]: rmse = 0
for i in range(len(predictions)):

    rmse += (list(Y_test)[i] - predictions[i]) ** 2

    rmse_out = np.sqrt(rmse / len(predictions))

print("RMSE:", rmse_out) # Pretty low
```

RMSE: 3.1906255142759106

```
[142]: model.intercept_
```

```
[142]: 3.355518994093951
```

```
[143]: print("Coefficient Values")
print("overall_sent_score:", model.coef_[0])
print("controversiality:", model.coef_[1])
print("body_length:", model.coef_[2])
```

Coefficient Values

```
overall_sent_score: -0.7711052668422775
controversiality: -2.2835054224168094
body_length: -0.007888327217906
```

6.2 Run general linear regression without splitting to get coefficients

```
[144]: import statsmodels.formula.api as sm # import statsmodels
```

- Regress the score of the comment on the length, it's overall sentiment score, and its controversiality.

```
[145]: results = sm.ols("score ~ body_length + overall_sent_score + controversiality",
                        data = politics_out).fit()
```

- Low R-Squared Value -> statistical significance for all dependent variables.

```
[290]: print(results.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  score    R-squared:                  0.022
Model:                            OLS    Adj. R-squared:             0.022
Method:                 Least Squares    F-statistic:                 158.9
Date:                Thu, 23 Apr 2020    Prob (F-statistic):          6.71e-102
Time:                  10:57:39    Log-Likelihood:              -55394.
No. Observations:                21461    AIC:                        1.108e+05
Df Residuals:                    21457    BIC:                        1.108e+05
Df Model:                          3
Covariance Type:                nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
Intercept                3.2959      0.105     31.296      0.000      3.090
3.502
controversiality[T.1]    -2.3192      0.118    -19.576      0.000     -2.551
-2.087
body_length              -0.0083      0.001     -9.265      0.000     -0.010
-0.007
overall_sent_score       -0.6039      0.263     -2.300      0.021     -1.119
-0.089
=====
Omnibus:                 3572.816    Durbin-Watson:              1.973
Prob(Omnibus):            0.000    Jarque-Bera (JB):           6124.309
Skew:                     1.094    Prob(JB):                    0.00
Kurtosis:                 4.436    Cond. No.                    396.

```

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

- Parameter Results: very similar to train/test split coefficients (6.1).

```
[147]: results.params # very very similar to train/test split coefficients
```

```
[147]: Intercept          3.295924
controversiality[T.1]  -2.319194
body_length           -0.008282
overall_sent_score    -0.603923
dtype: float64
```

6.3 Train regression model on TF-IDF matrix

- Re-define TF-IDF matrix

```
[333]: tfidf_vectorizer = TfidfVectorizer(max_df = 0.90, min_df = 10,
                                       ngram_range=(1, 1), stop_words = "english")
X = tfidf_vectorizer.fit_transform(corpus)

X = X.toarray()
tfidf_df = pd.DataFrame(X)
```

- Add score variable

```
[149]: tfidf_df = tfidf_df.join(politics_data[["score", "body_length",
↪ "overall_sent_score", "controversiality"]])
```

```
[150]: tfidf_df.head()
```

```
[150]:
```

	0	1	2	3	4	5	6	7	8	9	...	4100	4101	4102	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	

	4103	4104	4105	score	body_length	overall_sent_score	controversiality
0	0.0	0.0	0.0	1	9	0.468815	0
1	0.0	0.0	0.0	12	10	0.399522	0
2	0.0	0.0	0.0	1	7	0.389751	0
3	0.0	0.0	0.0	-6	6	0.390280	0
4	0.0	0.0	0.0	17	6	0.345299	0

[5 rows x 4110 columns]

- Remove outliers, subset y variable, remove score variable.

```
[151]: # Find IQR
Q1 = tfidf_df["score"].quantile(0.25)
Q3 = tfidf_df["score"].quantile(0.75)
IQR = Q3 - Q1

# Remove outliers
tfidf_out = tfidf_df[tfidf_df["score"] < (Q3 + 1.5 * IQR)]
tfidf_out = tfidf_out[tfidf_out["score"] > (Q1 - 1.5 * IQR)]

# subset y variable
y2 = tfidf_out.score

# drop score
tfidf_out = tfidf_out.drop(columns = ["score"])

# Reset Index
tfidf_out = tfidf_out.reset_index()
```

```
[152]: tfidf_out.head()
```

```
[152]:
```

	index	0	1	2	3	4	5	6	7	8	...	4099	4100	4101	\
0	0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
1	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
2	2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
3	3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
4	5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	

	4102	4103	4104	4105	body_length	overall_sent_score	controversiality
0	0.0	0.0	0.0	0.0	9	0.468815	0
1	0.0	0.0	0.0	0.0	10	0.399522	0
2	0.0	0.0	0.0	0.0	7	0.389751	0
3	0.0	0.0	0.0	0.0	6	0.390280	0
4	0.0	0.0	0.0	0.0	12	0.432123	0

[5 rows x 4110 columns]

- Train/test split

```
[153]: # split train and test data
X2_train, X2_test, Y2_train, Y2_test = train_test_split(tfidf_out, y2,
                                                         test_size = 0.25,
                                                         random_state = 1)
```

- Fit a new model

```
[154]: lm = linear_model.LinearRegression()
model2 = lm.fit(X2_train, Y2_train)
predictions2 = lm.predict(X2_test)
```

- **Produce plots of yhat vs. ytrue:** According to the plots below, the difference between y_{pred} and y_{true} is large sometimes but it is pretty close most of the time.

Reference: <https://bit.ly/2RT2SRe>


```
[286]: # create a df
df2 = pd.DataFrame({'Actual': Y2_test, 'Predicted': predictions2})
df2 = df.reset_index(drop = True)

# plotting
plt.figure(figsize = (16, 4))
df2.plot(kind = 'bar',figsize = (16,5))
plt.grid(which = 'major', linestyle = '-', linewidth = '0.5', color = 'green')
plt.grid(which = 'minor', linestyle = ':', linewidth = '0.5', color = 'black')
plt.xlim(1400,1430)
plt.title("Comparison between y_true & y_pred", fontsize = 18)
plt.ylabel("Value", fontsize = 14)
plt.xlabel("Order of Entries", fontsize = 14)

# take a look at part of the data
plt.xlim(1500,1530)
```

[286]: (1500, 1530)


<Figure size 1152x288 with 0 Axes>



```
[277]: plt.scatter(Y2_test, predictions2)
plt.title("Scatter Plot of y_true v.s. y_pred")
plt.xlabel("True Values")
```

```
plt.ylabel("Predictions")
```

```
[277]: Text(0, 0.5, 'Predictions')
```



- Summary Stats... R-squared not accurate (and negative \rightarrow not accurate \rightarrow see help page) but RMSE is low.

```
[157]: print("R-squared:", model2.score(X2_test, Y2_test))
```

R-squared: -0.3442531730282339

```
[158]: rmse = 0
for i in range(len(predictions2)):

    rmse += (list(Y2_test)[i] - predictions2[i]) ** 2

    rmse_out = np.sqrt(rmse / len(predictions2))

print("RMSE:", rmse_out)
```

RMSE: 3.745658641619775

- Make a dataframe from the coefficient outputs.


```
[159]: betas = pd.DataFrame(model2.coef_[1:])
      betas.columns = ["Betas"]
```

```
[184]: betas.head(10)
```

```
[184]:      Betas
0 -0.200438
1  1.276361
2  0.062117
3  3.030117
4 -0.700279
5 -0.813882
6  1.415927
7  5.907177
8 -0.846506
9  8.524779
```

- Sort term dictionary by the index that they were in for the TF-IDF matrix.

Reference: <https://stackoverflow.com/questions/613183/how-do-i-sort-a-dictionary-by-value>

```
[161]: terms_dict = {k: v for k, v in sorted(tfidf_vectorizer.vocabulary_.items(),
                                           key = lambda item: item[1])}
      terms = pd.DataFrame(list(terms_dict.keys()))
```

```
[183]: terms.head(10)
```

```
[183]:      0
0  abandoned
1   ability
2    able
3  abortion
4   abrams
5  absolute
6 absolutely
7   absurd
8   abuse
9   abused
```

- Join the dictionaries.

```
[163]: coefficients = betas.join(terms)
      coefficients.columns = ["Betas", "Feature"]
```

- See which 1-gram words had highest/lowest coefficients. There seems to be no discernable pattern.

```
[164]: coefficients.sort_values(by = 'Betas', ascending = False).head(7)
```

```
[164]:
```

	Betas	Feature
2655	78.391272	performed
2463	62.098451	notability
3416	62.098451	soundclips
2693	35.607651	placed
2020	26.326307	keywords
116	20.087543	alleviate
556	17.381032	cheated

```
[165]: coefficients.sort_values(by = "Betas", ascending = True).head(7)
```

```
[165]:
```

	Betas	Feature
3558	-121.007879	subredditmessagecomposeto
1260	-32.885438	exclusive
681	-27.632830	completing
249	-22.158693	authorize
1096	-19.887569	drama
1415	-17.731639	flooding
1059	-17.497666	division

7 Topic Modeling

- Create and apply a function trimming too long/possibly non-word. Remove words which contain more than 20 characters.

```
[ ]: # write a function
def vocab_trimmer(docs):

    words = docs.split()
    word_trim = [w for w in words if len(w) < 20 ]

    return(word_trim)

# apply the function for the body variable in politics_data
politics_data["body"] = politics_data["body"].apply(vocab_trimmer)

# make a corpus
corpus = list(politics_data["body"])

# briefly print the corpus
corpus[0:7]
```

7.1 TF-IDF vectorizing

```
[166]: tfidf_vectorizer = TfidfVectorizer(max_df = 0.90, min_df = 10,
                                         ngram_range=(1, 2), stop_words = "english")
X = tfidf_vectorizer.fit_transform(corpus)
tfidf_vocab = tfidf_vectorizer.get_feature_names()  # vocabulary list (terms)

X.toarray()  # term-freq matrix_At
```

```
[166]: array([[0., 0., 0., ..., 0., 0., 0.],
              [0., 0., 0., ..., 0., 0., 0.],
              [0., 0., 0., ..., 0., 0., 0.],
              ...,
              [0., 0., 0., ..., 0., 0., 0.],
              [0., 0., 0., ..., 0., 0., 0.],
              [0., 0., 0., ..., 0., 0., 0.]])
```

```
[167]: np.shape(X.toarray())  # 25,000 documents * 5,920 vocabulary
```

```
[167]: (25000, 5877)
```

7.2 NMF on TF-IDF matrix

```
[168]: ### NMF

# function for returning output
def print_top_words(model, vocab_list, n_top_words):
    for topic_idx, topic in enumerate(model.components_):
        message = "Topic #%d: " % topic_idx
        message += " ".join([vocab_list[i]
                             for i in topic.argsort()[::-n_top_words - 1:-1]])
        output = print(message)
    return(output)

# NMF results - 10 topics with top 20 words
nmf = NMF(n_components = 10).fit(X)
print_top_words(nmf, tfidf_vocab, n_top_words = 20)
```

Topic #0: violation rule automatically subredditmessagecomposeto question
performed automatically action performed bot action contact moderator moderator
subredditmessagecomposeto subredditmessagecomposeto automatically contact
question concern performed report bot ban comment harm rule rule violation
wishing deathphysical advocating wishing comment violation
Topic #1: people right republican know time want vote thing going need good make
way democrat really point sure year mean let
Topic #2: submission removal regarding removal regarding question removal
submission thank removed megathread feel free removal feel moderator regarding
question removal free message thank participating question regarding message

moderator participating message free

Topic #3: barr mueller said letter summary congress medium doj mueller said report testify thought inaccurate official memo mueller letter mueller report lied testimony coverage

Topic #4: like look look like sound sound like feel like feel barr look like trump like barr guy people like post lol lot act act like eye thing like thing

Topic #5: trump president supporter trump supporter biden donald obstruction donald trump crime justice evidence russia election investigation campaign like trump russian impeachment collusion trump campaign

Topic #6: public letter investigation department context summary counsel special special counsel substance nature substance context nature nature work capture conclusion capture context released march fully

Topic #7: read report say read report obstruction mueller report collusion article president report say justice evidence page lol mueller read article crime thing redacted saying

Topic #8: graham lindsey lindsey graham shit fucking fuck email hillary fucking idiot idiot lindsay clinton lindsay graham lol talking piece shit piece trump fucking hearing holy

Topic #9: think really think barr got people think think trump make guy wrong think going happen going tell really think job think mueller actually reason think mean lol

8 KMeans Clustering

- Import kmeans module

```
[302]: from sklearn.cluster import KMeans
```

- Subsampled data to test on

```
[303]: test = politics_data.sample(5000)
test_list = list(test["body"])
```

- Get TF-IDF and H Matrix (for clustering and topic output).

```
[304]: X = tfidf_vectorizer.fit_transform(corpus)
H = nmf.fit_transform(X)
vocab = tfidf_vectorizer.get_feature_names()
```

- Write a function to get top 10 words which are the nearest to the centroids.

Reference: <http://brandonrose.org/clustering#Tf-idf-and-document-similarity>

```
[305]: ## Top n words for each cluster
def top_words_KM(cluster, n_top_words, vocab_list):

    # get centroids and arrange the indices of centroids
    order_centroids_index = cluster.cluster_centers_.argsort()[:, :-1]
```

```

    for i in range(cluster.n_clusters):
        print_wds = [vocab_list[index] for index in order_centroids_index[i,:
→n_top_words]]
        output = print("Cluster {}: {}".format(i, ' '.join(print_wds)))

    return(output)

```

```

[306]: KM = KMeans(n_clusters = 10, random_state = 1, max_iter = 100)
kmeans_label = KM.fit_predict(X).tolist()
top_words_KM(cluster = KM, n_top_words = 15, vocab_list = vocab)

```

Cluster 0: right, graham, shit, fucking, lindsey, lindsey graham, trump, fuck, like, thing, piece shit, know, piece, going, email

Cluster 1: submission, removal, regarding removal, regarding, removal submission, question, thank, removed, megathread, removal feel, free message, moderator regarding, question removal, thank participating, question regarding

Cluster 2: people, republican, vote, trump, like, think, want, democrat, thing, know, party, need, right, make, american

Cluster 3: violation, rule, shill troll, attack idea, general courteous, personal insult, troll accusation, accusation hate, advocating wishing, result permanent, permanent ban, subreddit civil, ban comment, violation result, speech advocating

Cluster 4: mueller, barr, letter, said, doj, report, summary, congress, mueller said, medium, mueller letter, say, investigation, testify, know

Cluster 5: trump, like, think, know, say, going, need, want, make, president, thing, way, really, mean, got

Cluster 6: barr, trump, like, summary, letter, report, congress, mueller, going, think, know, tomorrow, lie, said, testimony

Cluster 7: good, good thing, thing, luck, good luck, trump, guy, make, like, know, think, bad, right, point, really

Cluster 8: time, trump, like, think, going, people, actually, right, thing, know, year, need, barr, long, vote

Cluster 9: report, read, mueller, summary, obstruction, read report, mueller report, public, counsel, letter, special counsel, special, context, investigation, conclusion

- Get the counts of the cluster assignments.

```

[307]: from collections import Counter
Counter(kmeans_label) # one cluster has most of the comments (seems like the_
→most general too)

```

```

[307]: Counter({7: 614,
               5: 14756,
               2: 2537,
               4: 1479,
               3: 469,
               9: 1070,


```

```
6: 1399,  
1: 231,  
8: 840,  
0: 1605})
```

- Plot a histogram of the cluster assignments.

```
[325]: plt.hist(kmeans_label, color = "coral")  
plt.xlabel("Clusters")  
plt.ylabel("# of Observations in Each Cluster")  
plt.title("Cluster Scoring Distribution")
```


```
[325]: Text(0.5, 1.0, 'Cluster Scoring Distribution')
```



- Plot a density plot of the cluster assignments.

```
[309]: sns.distplot(kmeans_label, hist = False, color= "blue", bins = 1)  
plt.title("Density Plot of Cluster Distribution")  
plt.xlabel("Clusters")  
plt.ylabel("Proportion in Each Cluster")
```

```
[309]: Text(0, 0.5, 'Proportion in Each Cluster')
```



- Check text of comments to see if they are similar.

```
[320]: my_list = []
for i in range(len(kmeans_label)):
    if kmeans_label[i] == 5:
        my_list.append(i)

print(my_list[0:10])
print(my_list[145:155])
print(my_list[-10:-1])
```

```
[1, 4, 5, 7, 9, 11, 12, 14, 15, 20]
[242, 243, 246, 247, 248, 249, 252, 253, 254, 259]
[24982, 24984, 24985, 24987, 24990, 24992, 24993, 24994, 24996]
```

- Pick a few index numbers from directly above and replace in code below to see similar output. The clusters seem to line up.

```
[324]: print("1.", politics_data["body"][5])

print("2", politics_data["body"][243])

print("3",politics_data["body"][24992])
```

1. everything power make sure biden get nomination fall line doe primary

coronation
 2 nope nonstate function dignitary conducting diplomacy actual state function
 govt funded kind event actually considered government business reimbursed thus
 trump pr move buying hamberders shutdown save government intended save him
 billionaire money
 3 well corporation democrat pocket case trump make cut

9 Classification

9.1 Two classes - sports subreddits

- We wanted to pick subreddits that were similar enough to be tough to train a model on but distinct enough to tell differences. Sports themes subreddits seemed to fit the model.

```
[197]: # Re-load original data
data = data
```

- Import packages

```
[198]: # classifying modules
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import RidgeClassifier

# sklearn
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```

- Subset data

```
[199]: subs = ["nba", "nfl"]
test = data[data.subreddit.isin(subs)]
```

- Clean data

```
[200]: test["body"] = test["body"].apply(textCleaner)
test["body"] = test["body"].apply(textNormalization)
test["body"]
```

```
[200]: 7                report water wet
12                michael jackson popcorn gif
23    the reason anyone know jared dudley ben simmon...
31    would say year old buddy hield part king young...
55                oh know going comment lol

...
919444    mean put weight letter lawyer nfl gtfinally m ...
919445                rip fedora man really getting good
919449    andy reid marv levy suffer lack super bowl alt...
```



```

919456    none change fact succeeded nfl point star play...
919486    mean reason season great last year oline stren...
Name: body, Length: 50000, dtype: object

```

- Define and run classifier function

```

[201]: def isNBA(string):
        x = 0
        if string == "nba":
            x += 1
        return(x)

test["category"] = test.subreddit.apply(isNBA)

```

```

[203]: test.head()

```

```

[203]:      subreddit                                     body \
7         nba                                     report water wet
12        nba                                michael jackson popcorn gif
23        nba  the reason anyone know jared dudley ben simmon...
31        nba  would say year old buddy hield part king young...
55        nba                                oh know going comment lol

      controversiality  score  category
7                    0      9         1
12                   0      1         1
23                   0     -5         1
31                   1     -2         1
55                   0      1         1

```

- Make corpus for vectorization and y_true tags.

```

[204]: text = list(test["body"])
        classifier = list(test["category"])

```

- Split and vectorize the data

```

[205]: data_train, data_test = train_test_split(text, train_size = 0.8, random_state = 5)
        y_train, y_actual = train_test_split(classifier, train_size = 0.8, random_state = 5)

```

```

[206]: X_train = tfidf_vectorizer.fit_transform(data_train)
        H_train = nmf.fit_transform(X_train)

```

- Fit Knn classifier

```
[207]: neigh = KNeighborsClassifier(n_neighbors = 4)
       neigh.fit(H_train, y_train)
```

```
[207]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                           metric_params=None, n_jobs=None, n_neighbors=4, p=2,
                           weights='uniform')
```

- Generate predicted classes

```
[208]: X_test = tfidf_vectorizer.transform(data_test) # not fit
       H_test = nmf.transform(X_test) # not fit
```

```
[209]: y_pred = neigh.predict(H_test)
```

- Get classifier scores. This model is overall a pretty accurate model.

```
[210]: print("F1 Score:", f1_score(y_actual, y_pred))
       print("Precision Score:", precision_score(y_actual, y_pred))
       print("Recall Score:", recall_score(y_actual, y_pred))
```

```
F1 Score: 0.6695862625331336
Precision Score: 0.7911220043572985
Recall Score: 0.5804195804195804
```

```
[211]: target_names = ["NBA", "NFL"]
       print(classification_report(y_actual, y_pred, target_names = target_names))
```

	precision	recall	f1-score	support
NBA	0.67	0.85	0.75	4995
NFL	0.79	0.58	0.67	5005
accuracy			0.71	10000
macro avg	0.73	0.71	0.71	10000
weighted avg	0.73	0.71	0.71	10000

9.2 Four classes - More sports

- According to the classification results below, it is tougher to train more classes.

```
[212]: data = data
```

```
[213]: subs = ["nba", "nfl", "hockey", "soccer"]
       test = data[data.subreddit.isin(subs)]
```

```
[214]: test["body"] = test["body"].apply(textCleaner)
       test["body"] = test["body"].apply(textNormalization)
```

```
test["body"]
```

```
[214]: 7          report water wet
      12          michael jackson popcorn gif
      21          google chrome suck donkey dick
      23  the reason anyone know jared dudley ben simmon...
      31  would say year old buddy hield part king young...
      ...
919444  mean put weight letter lawyer nfl gtfinally m ...
919445          rip fedora man really getting good
919449  andy reid marv levy suffer lack super bowl alt...
919456  none change fact succeeded nfl point star play...
919486  mean reason season great last year oline stren...
Name: body, Length: 100000, dtype: object
```

```
[215]: def isNBA(string):
      x = 0
      if string == "nba":
          x = 0
      elif string == "nfl":
          x = 1
      elif string == "hockey":
          x = 2
      elif string == "soccer":
          x = 3
      return(x)

test["category"] = test.subreddit.apply(isNBA)
```

```
[216]: test
```

```
[216]:      subreddit      body \
7          nba          report water wet
12         nba          michael jackson popcorn gif
21        hockey          google chrome suck donkey dick
23         nba  the reason anyone know jared dudley ben simmon...
31         nba  would say year old buddy hield part king young...
...
919444      nfl  mean put weight letter lawyer nfl gtfinally m ...
919445      nfl          rip fedora man really getting good
919449      nfl  andy reid marv levy suffer lack super bowl alt...
919456      nfl  none change fact succeeded nfl point star play...
919486      nfl  mean reason season great last year oline stren...

      controversiality  score  category
7                      0      9        0
12                     0      1        0
```

21	0	2	2
23	0	-5	0
31	1	-2	0
...
919444	0	0	1
919445	0	4	1
919449	0	7	1
919456	0	4	1
919486	0	11	1

[100000 rows x 5 columns]

```
[217]: text = list(test["body"])
classifier = list(test["category"])
```

```
[218]: data_train, data_test = train_test_split(text, train_size = 0.8, random_state = 5)
y_train, y_actual = train_test_split(classifier, train_size = 0.8, random_state = 5)
```

```
[219]: X_train = tfidf_vectorizer.fit_transform(data_train)
H_train = nmf.fit_transform(X_train)
```

```
[220]: neigh = KNeighborsClassifier(n_neighbors = 4)
neigh.fit(H_train, y_train)
```

```
[220]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=4, p=2,
weights='uniform')
```

```
[221]: X_test = tfidf_vectorizer.transform(data_test) # not fit
H_test = nmf.transform(X_test) # not fit
```

```
[222]: y_pred = neigh.predict(H_test)
```

```
[223]: f1_score(y_actual, y_pred, average="weighted")
```

```
[223]: 0.3836387209725444
```

```
[224]: precision_score(y_actual, y_pred, average="weighted")
```

```
[224]: 0.3886078659789782
```

```
[225]: recall_score(y_actual, y_pred, average = "weighted")
```

```
[225]: 0.38765
```

- The output gives a less accurate model, but still not that bad.

```
[226]: target_names = ["NBA", "NFL", "Hockey", "Soccer"]
print(classification_report(y_actual, y_pred, target_names = target_names))
```

	precision	recall	f1-score	support
NBA	0.41	0.51	0.46	5003
NFL	0.36	0.42	0.39	4921
Hockey	0.35	0.29	0.32	5011
Soccer	0.43	0.33	0.37	5065
accuracy			0.39	20000
macro avg	0.39	0.39	0.38	20000
weighted avg	0.39	0.39	0.38	20000

10 Appendix

10.1 Bindary Logistic Regression

- As we can see from the counts & graph below, the percentage of controversiality is 96.776, and the percentage of no controversiality is 3.22399. The controversiality is unbalanced.

Reference <https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8>

```
[243]: politics_data["controversiality"].value_counts()[1]
```

```
[243]: 806
```

```
[244]: # see the counts of controversiality
print("Value counts of controversiality:")
print(politics_data["controversiality"].value_counts())
```

```
Value counts of controversiality:
0    24194
1      806
Name: controversiality, dtype: int64
```

- If we look at the mean of score and sentiment scores groupby controversiality, the average score of up reddits minus down reddits has discrepancy (2.9 v.s. 0.58). However, the mean in all sentiment scores as well as overall sentiment score are all pretty close in between. Even the overall score has 0.01 higher in controversiality then no controversiality.

```
[245]: politics_out.groupby("controversiality").mean()
```

```
[245]:          index    score  body_length  Positive_Score \
controversiality
```

0	12485.650826	2.908529	18.535642	0.108551
1	12918.679470	0.582781	18.998675	0.108177

	Negative_Score	net_score	slang_score	anew_score	\
controversiality					
0	0.106273	0.501139	0.472363	0.070580	
1	0.106746	0.500716	0.472868	0.071179	

	vader_score	overall_sent_score
controversiality		
0	0.505058	0.387285
1	0.522387	0.391787

- Implement the model. The p-values are small (< 0.05).

```
[255]: import statsmodels.api as sm
# subset y variable
x = politics_out[["overall_sent_score", "score"]]
y = politics_out.controversiality

# Logit model
logit_model = sm.Logit(y, x)
result = logit_model.fit()

# ----- print result -----
# p-value < 0.05 (both score & overall_sent_score)
result.summary2()
```

Optimization terminated successfully.
Current function value: 0.147881
Iterations 8

```
[255]: <class 'statsmodels.iolib.summary2.Summary'>
"""
```

```

Results: Logit
=====
Model:                Logit                Pseudo R-squared: 0.029
Dependent Variable:    controversiality AIC:                6351.3343
Date:                 2020-04-23 04:58 BIC:                6367.2822
No. Observations:     21461                Log-Likelihood:    -3173.7
Df Model:              1                    LL-Null:         -3268.8
Df Residuals:          21459                LLR p-value:      2.9149e-43
Converged:             1.0000                Scale:           1.0000
No. Iterations:        8.0000

-----
              Coef.   Std.Err.    z      P>|z|    [0.025   0.975]
-----
```

```
overall_sent_score -7.1376    0.1050 -67.9966 0.0000 -7.3433 -6.9319
score              -0.3505    0.0152 -23.0740 0.0000 -0.3803 -0.3208
=====
```

```
"""
```

- Continually try on Logistic Regression Model Fitting.

```
[256]: from sklearn import linear_model
from sklearn import metrics

X_train, X_test, y_train, y_test = train_test_split(x, y,
                                                    test_size = 0.25,
                                                    random_state = 1)
lm = linear_model.LogisticRegression()
model = lm.fit(X_train, Y_train)
predictions = lm.predict(X_test)

# ----- print result -----
print("Intercept:", model.intercept_)
print("RMSE:", np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
Intercept: [ -7.69446842 -5.76404902 -4.89978361 -5.11096805 -3.74612598
 -3.54537763 -2.3821199    0.60505876 -1.40963169 -2.33395375
 -3.11511531 -3.5133533   -3.98212319 -4.42548451 -5.30640502
 -6.13031765 -6.5206707   -7.84194941 -9.34061325 -13.52020389]
RMSE: 2.80211437127095
```

- **Accuracy (R-squared):** The accuracy of logistic regression classifier on test set is 0.96 (very high).

```
[257]: print("R-squared:", model.score(X_test, y_test))
```

```
R-squared: 0.03410361535594484
```

- **Confusion Matrix:** The result gives that we have [5168 + 0] correct predictions and [198 + 0] incorrect predictions.

```
[258]: from sklearn.metrics import confusion_matrix

confusion_matrix = confusion_matrix(y_test, predictions)
print(confusion_matrix)
```

```
[[ 0 3963  501  306   77   49  110   25   17  120]
 [ 0  183    4    5    1    0    4    0    0    1]
 [ 0    0    0    0    0    0    0    0    0    0]
 [ 0    0    0    0    0    0    0    0    0    0]
 [ 0    0    0    0    0    0    0    0    0    0]
 [ 0    0    0    0    0    0    0    0    0    0]
```

```
[ 0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0]]
```

Precision, recall, F-measure and support

The precision not to label classifier positive if it's negative. The recall helps to find all positive samples, in this data result, we have most non-controversiality samples. Lastly, the f1-score gives an overall mean to evaluate precision and recall. The support returns the number of occurrences of each class in `y_test`.

According to the report below and the unbalanced controversiality rate, the precision and recall are both high (precision = 96%, recall = 100%), though the f1-score = 98% is almost the best. I believe this data is overfitting. Of the entire dataset, 98% reviews are non-controversiality.

```
[253]: from sklearn.metrics import classification_report

print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	5168
1	0.04	0.92	0.08	198
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
5	0.00	0.00	0.00	0
6	0.00	0.00	0.00	0
7	0.00	0.00	0.00	0
9	0.00	0.00	0.00	0
10	0.00	0.00	0.00	0
13	0.00	0.00	0.00	0
accuracy			0.03	5366
macro avg	0.00	0.09	0.01	5366
weighted avg	0.00	0.03	0.00	5366

ROC Curve

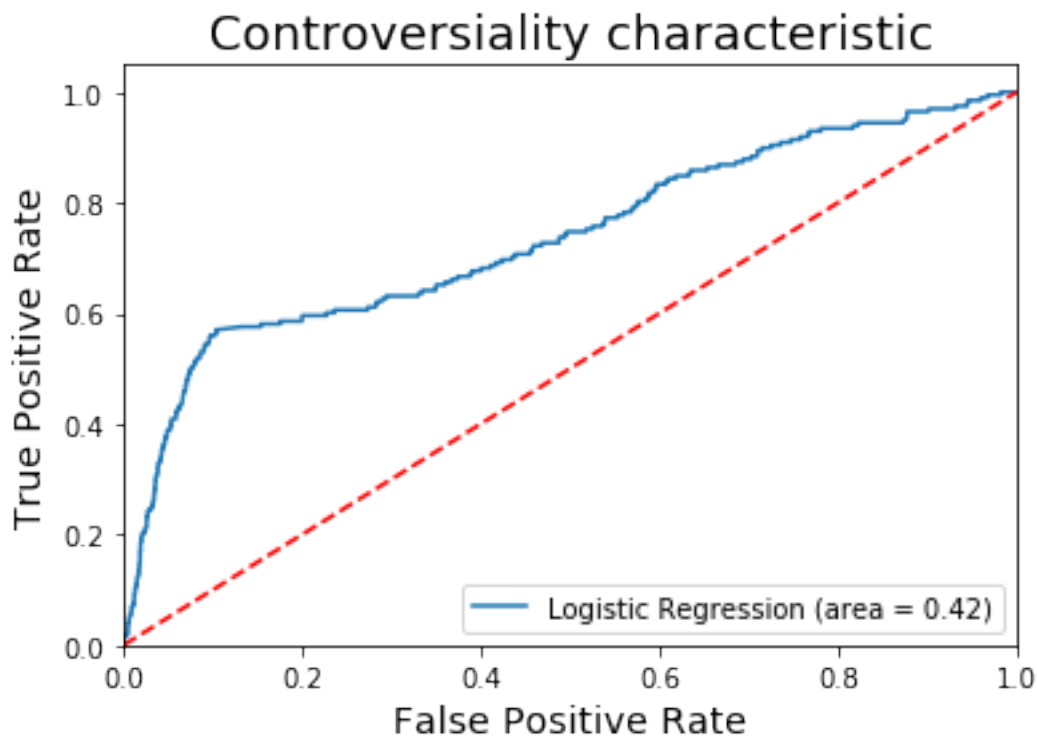
```
[254]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve

logit_roc_auc = roc_auc_score(y_test, predictions)
fpr, tpr, thresholds = roc_curve(y_test, lm.predict_proba(X_test)[: ,1])

plt.figure()
plt.plot(fpr, tpr, label = 'Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
```



```
plt.ylim([0.0, 1.05])
plt.title('Controversiality characteristic', fontsize = 18)
plt.xlabel('False Positive Rate', fontsize = 14)
plt.ylabel('True Positive Rate', fontsize = 14)
plt.legend(loc = "lower right")
plt.savefig('Log_ROC')
plt.show()
```



10.1.1 /Additional Work/

- We tried to create a visual graph of the clusters; however, this proved to be too difficult. To do this, we found that we needed to convert each document to a 2D form so it could be plotted on a flat plane. The computing power to do this on the entire dataset and produce coherent output was very difficult. The output graph below was with 5,000 of the 25,000 observations, and the online source we were trying to work through was difficult to apply to our project. We ultimately felt it was best to leave it out.

10.2 Multidimensional Scaling

- Using multidimensional scaling to convert the dist matrix into a 2-dimensional array.

Reference: <http://brandonrose.org/clustering#Tf-idf-and-document-similarity>*

- Get cosine distance

```
[227]: from sklearn.metrics.pairwise import cosine_similarity
dist = 1 - cosine_similarity(X)
```

- Import packages

```
[228]: import matplotlib as mpl
from sklearn.manifold import MDS
```

```
[229]: # convert two components as we're plotting points in a two-dimensional plane
# "precomputed" because we provide a distance matrix
MDS()
mds = MDS(n_components = 2, dissimilarity = "precomputed", random_state = 1,
↳max_iter = 100)
```

- The transform function below for mds() takes a very very long while running through it.

```
[347]: # [!] takes a while
pos = mds.fit_transform(dist) # shape (n_components, n_samples)
xs, ys = pos[:, 0], pos[:, 1]
```

```
[306]: # create a dataframe
cluster_df = pd.DataFrame(dict(x = xs, y = ys,
                                label = kmeans_label,
                                score = list(test["score"]))) # should be
↳title = titles
# group by clusters
groups = cluster_df.groupby('label')
```

```
[230]: # set up colors per clusters using a dict
cluster_colors = {0: 'red', 1: 'blue', 2: 'green', 3: 'purple', 4: 'orange', 5:
↳'brown'}

#set up cluster names using a dict
cluster_names = {0: 'republican, Trump, democrat',
                  1: 'justice, Mueller report, congress',
                  2: 'violation, harm',
                  3: 'thoughts, feelings',
                  4: 'say, Trump, action',
                  5: 'submission, removal'}
```

```
[308]: # set output plot size
fig, ax = plt.subplots(figsize = (17, 9)) # set size
ax.margins(0.05)

# iterate through groups to layer the plot
for name, group in groups:
    ax.plot(group.x, group.y, marker = 'o', linestyle = '', ms = 12,
```

```

        label = cluster_names[name], color = cluster_colors[name],
        mec = 'none')
ax.set_aspect('auto')
ax.tick_params(\
    axis = 'x',          # changes apply to the x-axis
    which = 'both',      # both major and minor ticks are affected
    bottom = 'off',      # ticks along the bottom edge are off
    top = 'off',         # ticks along the top edge are off
    labelbottom = 'off')
ax.tick_params(\
    axis = 'y',          # changes apply to the y-axis
    which = 'both',      # both major and minor ticks are affected
    left = 'off',        # ticks along the bottom edge are off
    top = 'off',         # ticks along the top edge are off
    labelleft = 'off')

ax.legend(numpoints = 1)

# add labels
for i in range(len(cluster_df)):
    ax.text(cluster_df.iloc[i]['x'],
            cluster_df.iloc[i]['y'],
            cluster_df.iloc[i]['score'], size = 9)

plt.show() #show the plot

```

